

OS/161 – NEW MEMORY MANAGER

Peter Coffman

Ryan Basile

OVERVIEW

- OS/161 is a simple operating system
 - Many of its features are unimplemented, and it's up to you to implement them
- Memory management is especially primitive
 - Memory is never freed
 - Memory is always allocated in contiguous chunks

PROBLEM 1

- OS/161 does not properly deallocate memory when threads exit
 - This caused the test programs for lab 4 to crash after running them 3 times

```
OS/161 kernel [? for menu]: sy2
Starting lock test...
scheduler: Dropping thread synctest.
panic: locktest: thread_fork failed: Out of memory
sys161: 429496729742640698 cycles (16830378 run, 429496729725810320 global-idle)
sys161:    cpu0: 16830378 kern, 0 user, 0 idle)
sys161: 1774 irqs 0 exns 0r/0w disk 12r/1247w console 0r/0w/1m emufs 0r/0w net
sys161: Elapsed real time: 6.696244 seconds (6.41399e+10 mhz)
sys161: Elapsed virtual time: 5.705627937 seconds (25 mhz)
os161user@os161user-VirtualBox:~/cs4300-os161/root$
```

```
OS/161 kernel [? for menu]: sy3
Starting CV test...
Threads should print out in reverse order.
scheduler: Dropping thread synctest.
panic: cvtest: thread_fork failed: Out of memory
sys161: 644245206064255968 cycles (746399833 run, 644245205317856135 global-idle)
sys161:    cpu0: 746399833 kern, 0 user, 0 idle)
sys161: 447869 irqs 0 exns 0r/0w disk 17r/4470w console 0r/0w/1m emufs 0r/0w net
sys161: Elapsed real time: 4466.591779 seconds (1.44236e+08 mhz)
sys161: Elapsed virtual time: 4466.570238737 seconds (25 mhz)
os161user@os161user-VirtualBox:~/cs4300-os161/root$
```

PROBLEM 2

- OS/161 eagerly allocates all memory of a process up front
 - Sparse access patterns are not supported

```
OS/161 kernel [? for menu]: p testbin/sparse
Running program testbin/sparse failed: Out of memory
Operation took 0.723996600 seconds
OS/161 kernel [? for menu]: █
```

OS ANALYSIS

- OS/16I comes with a very basic virtual memory manager, “dumbvm,” located in the dumbvm.c file
- It has two empty functions:

```
21 void
22 vm_bootstrap(void)
23 {
24     /* Do nothing. */
25 }
26
```

```
54 void
55 free_kpages(vaddr_t addr)
56 {
57     /* nothing */
58     (void)addr;
59 }
60
61
```

- These are called in main.c and kheap.c
- Also includes some basic address space functions

WHERE THEY'RE REFERENCED

MAIN.C

- Responsible for system boot, shutdown, and running the main kernel.
- It runs `vm_bootstrap()` during its `boot()` function

```
ram_bootstrap();  
scheduler_bootstrap();  
thread_bootstrap();  
vfs_bootstrap();  
dev_bootstrap();  
vm_bootstrap();  
kprintf_bootstrap();
```

KHEAP.C

- Subpage allocator
- `subpage_kfree()` function uses `free_kpages()` to free whole pages marked "free" given a virtual page address "prpage"
- `kfree()` also uses it
 - More details next...

KFREE.C

- Relates to the test programs sy2 and sy3:
 - Creates threads, which exit when completed (`thread_exit`)
 - Exited threads are marked as "zombie threads"
 - The `exorcise()` function tries to destroy the zombie threads with "`thread_destroy(z)`"
 - However, `thread_destroy()` uses `kfree()`, which doesn't work because `free_kpages()` is unimplemented!
 - Pages of memory get filled but are never freed
 - Eventually leads to "Out of memory" error

```
74 static
75 void
76 thread_destroy(struct thread *thread)
77 {
78     assert(thread != curthread);
79
80     // If you add things to the thread structure, be sure to dispose of
81     // them here or in thread_exit.
82
83     // These things are cleaned up in thread_exit.
84     assert(thread->t_vmspace==NULL);
85     assert(thread->t_cwd==NULL);
86
87     if (thread->t_stack) {
88         kfree(thread->t_stack);
89     }
90
91     kfree(thread->t_name);
92     kfree(thread);
93 }
```



```
565 void
566 kfree(void *ptr)
567 {
568     /*
569      * Try subpage first; if that fails, assume it's a big allocation.
570      */
571     if (ptr == NULL) {
572         return;
573     } else if (subpage_kfree(ptr)) {
574         assert((vaddr_t)ptr%PAGE_SIZE==0);
575         free_kpages((vaddr_t)ptr);
576     }
577 }
```

POSSIBLE SOLUTIONS

1. Implement the missing functions in dumbvm
 - Create a core map to keep track of pages (single table)
 - Modify getpages() to reference the core map
 - Add other functions to handle pages in the new core map
 - vm_bootstrap() to initialize the core map
2. Replace dumbvm with a more robust virtual memory manager

OUR SOLUTION

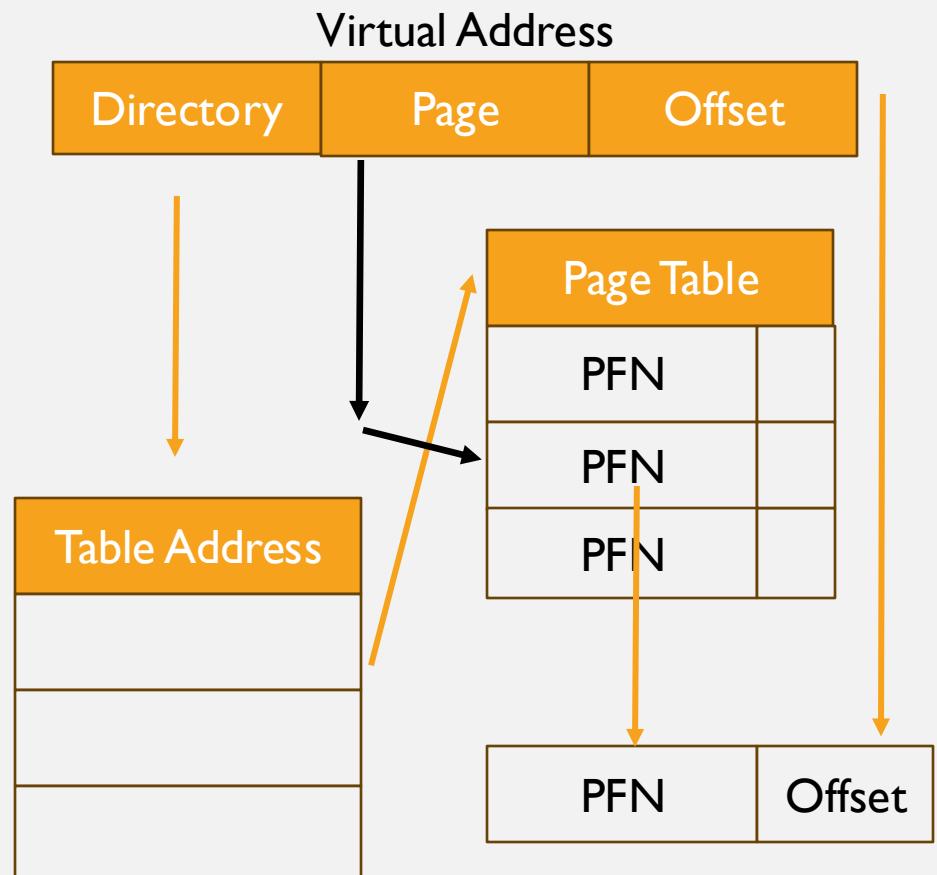
Replace dumbvm.c with vm.c

Uses a two-level page table

Saves memory

Add necessary functions,
modifying preexisting ones from
dumbvm

vm_bootstrap, free_kpages,
alloc_kpages, alloc_ppages...



FILE CHANGES

Many files had to be altered to allow two-level page tables.
A complete list is in the included documentation file.

- vm.c (new file)
 - Bootstrap function that allocates the physical memory on boot
 - Page allocation & free functions
 - For virtual pages & physical pages
- Page table nodes
- Fault handler
 - What happens if address is invalid?
- addrspace.c
 - Main structure for two levels
 - Defines virtual bases & number of pages for each level in an address space
 - Defines the stack

FILE CHANGES

- vm.h (new file)
 - Structure for page tables
 - Defines macro KVADDR_TOPADDR
 - Converts virtual address to physical address
 - Struct pt_node
 - Struct page_table
- addrspace.h
 - If compiled without dumbvm, redefine the variables:
 - vaddr_t, size_t, vaddr_t, page_table *pt
 - New function to check for legal virtual memory access: as_check_access
 - Prevent reading from subpage which isn't part of a page

TEST PROGRAMS

Two new programs were created to test the functionality of the new virtual memory system.

vm_t1.c

- A simple function that returns 0

vm_t2.c

- Creates a large array that's twice the size of physical memory
 - $2 \times 128 \times 4096$ ($2 \times \# \text{ of pages} \times \text{page size}$)
- Iterates through it, setting the value of the array at "i" to one
- Increments "i" to 10 times the page size (4096×10)
- Return 0 when completed

THE RESULTS

- Lab 4 test programs sy2 and sy3 no longer crash after three runs
- The sparse program returns 0
 - After a string of unknown syscall messages
- Both of our custom test programs return 0
 - No crashes or "out of memory" errors

```
OS/161 kernel [? for menu]: sy2
Starting lock test...
Lock test done.
Operation took 0.047091040 seconds
OS/161 kernel [? for menu]: sy2
Starting lock test...
Lock test done.
Operation took 0.047067200 seconds
OS/161 kernel [? for menu]: sy2
Starting lock test...
Lock test done.
Operation took 0.047067200 seconds
OS/161 kernel [? for menu]: █
```

```
OS/161 kernel [? for menu]: p /testbin/vm_t2
0
Operation took 0.975590040 seconds
OS/161 kernel [? for menu]:
```

```
Unknown syscall 6
0
Operation took 0.542768080 seconds
OS/161 kernel [? for menu]:
```